

# index

September 4, 2024

## 0.1 Overview

This project aims to develop a predictive model that can determine whether an arrest will be made during a Terry Stop based on various factors. The aim is to provide actionable insights that can help law enforcement agencies optimize their stop-and-search strategies, ensure public safety, and address concerns related to racial profiling and fairness in policing.

### 0.1.1 Business Understanding: Enhancing Law Enforcement Practices

Terry Stops, based on *Terry v. Ohio* (1968), are brief stops and searches by police based on reasonable suspicion. These stops raise concerns about fairness and potential bias. The business challenge is to predict the likelihood of an arrest during a Terry Stop to improve decision-making and address potential biases. The business understanding of the project is to enhance law enforcement practices by providing a platform for data-driven decision. This project benefits law enforcement by improving efficiency and fairness, policymakers by providing data-driven insights, and the public by fostering trust and reducing unjustified stops and arrests.

```
[109]: # importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

### 0.1.2 Data

The data has 61050 rows and 23 columns

```
[110]: # Load the data
terry_df = pd.read_csv('data/Terry_Stops.csv')
terry_df.head()
```

```
[110]:
```

	Subject	Age Group	Subject ID	GO / SC Num	Terry Stop ID \
0		1 - 17	-1	20170000315550	301338
1		26 - 35	-1	20170000115391	253596
2		18 - 25	-1	20160000036392	123863
3		36 - 45	33919122751	20210000055627	21770133324
4		18 - 25	-1	20160000003522	183376

	Stop Resolution	Weapon Type	Officer ID	Officer YOB	Officer Gender	\
0	Offense Report	NaN	5489	1964	M	
1	Offense Report	NaN	6403	1969	M	
2	Offense Report	NaN	7473	1981	M	
3	Field Contact	-	6711	1977	M	
4	Field Contact	NaN	7090	1981	F	

	Officer Race	...	Reported Time	\
0	White	...	00:10:00.0000000	
1	White	...	16:04:00.0000000	
2	White	...	02:26:00.0000000	
3	White	...	14:42:43.0000000	
4	White	...	15:49:00.0000000	

	Initial Call Type	\
0	NUISANCE	
1	THEFT (DOES NOT INCLUDE SHOPLIFT OR SVCS)	
2	SUSPICIOUS PERSON, VEHICLE, OR INCIDENT	
3	PROPERTY - DAMAGE	
4	-	

	Final Call Type	Call Type	\
0	--MISCHIEF OR NUISANCE - GENERAL	911	
1	--THEFT - ALL OTHER	TELEPHONE OTHER, NOT 911	
2	--SUSPICIOUS CIRCUM. - SUSPICIOUS PERSON	911	
3	--DISTURBANCE - OTHER	911	
4	-	-	

	Officer Squad	Arrest Flag	Frisk Flag	Precinct	\
0	NORTH PCT 3RD W - B/N RELIEF	N	N	North	
1	NORTH PCT 2ND W - LINCOLN - PLATOON 1	N	N	North	
2	NORTH PCT 3RD W - B/N RELIEF	N	N	North	
3	SOUTHWEST PCT 2ND W - WILLIAM - PLATOON 2	N	N	-	
4	SOUTHWEST PCT 2ND W - FRANK - PLATOON 2	N	Y	-	

	Sector	Beat
0	N	N2
1	L	L3
2	B	B3
3	-	-
4	-	-

[5 rows x 23 columns]

```
[111]: # Data information
        terry_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 61050 entries, 0 to 61049
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Subject Age Group                     61050 non-null  object
1   Subject ID                           61050 non-null  int64
2   GO / SC Num                          61050 non-null  int64
3   Terry Stop ID                        61050 non-null  int64
4   Stop Resolution                       61050 non-null  object
5   Weapon Type                          28485 non-null  object
6   Officer ID                           61050 non-null  object
7   Officer YOB                          61050 non-null  int64
8   Officer Gender                       61050 non-null  object
9   Officer Race                         61050 non-null  object
10  Subject Perceived Race                61050 non-null  object
11  Subject Perceived Gender              61050 non-null  object
12  Reported Date                        61050 non-null  object
13  Reported Time                        61050 non-null  object
14  Initial Call Type                    61050 non-null  object
15  Final Call Type                      61050 non-null  object
16  Call Type                           61050 non-null  object
17  Officer Squad                        60489 non-null  object
18  Arrest Flag                          61050 non-null  object
19  Frisk Flag                           61050 non-null  object
20  Precinct                             61050 non-null  object
21  Sector                               61050 non-null  object
22  Beat                                 61050 non-null  object
dtypes: int64(4), object(19)
memory usage: 10.7+ MB

```

```
[112]: terry_df.shape
```

```
[112]: (61050, 23)
```

## 0.2 Checking for Missing Values

The weapon type column has 32565 missing values and officer squad has 561 missing values.

```

[113]: # check for missing data in the dataset
missing_values = terry_df.isnull().sum()
missing_values

```

```

[113]: Subject Age Group      0
       Subject ID            0
       GO / SC Num           0
       Terry Stop ID         0
       Stop Resolution        0


```

Weapon Type	32565
Officer ID	0
Officer YOB	0
Officer Gender	0
Officer Race	0
Subject Perceived Race	0
Subject Perceived Gender	0
Reported Date	0
Reported Time	0
Initial Call Type	0
Final Call Type	0
Call Type	0
Officer Squad	561
Arrest Flag	0
Frisk Flag	0
Precinct	0
Sector	0
Beat	0
dtype:	int64

### 0.3 Analyze the Categorical Features

```
[114]: # List of categorical columns to analyze
categorical_cols = ['Subject Age Group', 'Stop Resolution', 'Subject Perceived_
↳ Race', 'Subject Perceived Gender']
```

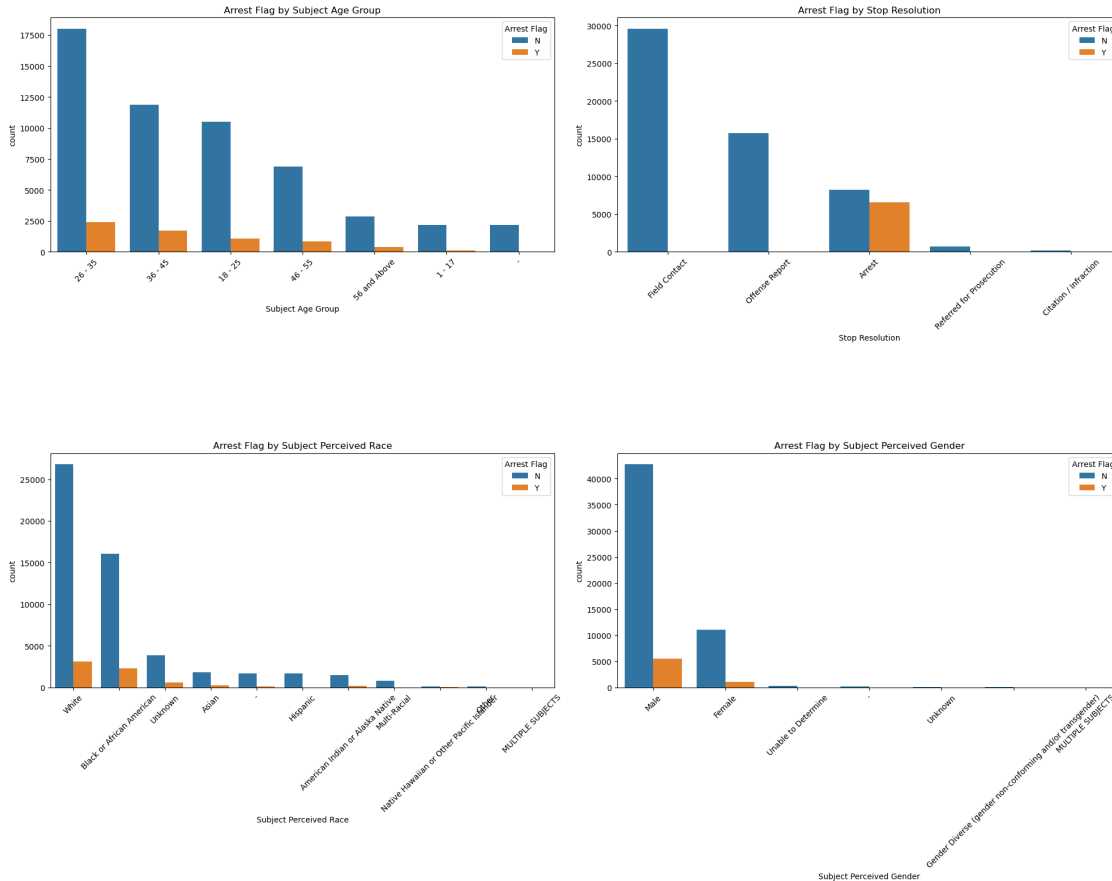
```
[115]: # Plot count plots for each categorical feature
plt.figure(figsize=(20, 12))
for i, col in enumerate(categorical_cols, 1):
    plt.subplot(3, 2, i)
    sns.countplot(x=col, data=terry_df, order=terry_df[col].value_counts().
↳ index)
    plt.title(f"Count Plot for {col}")
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## 0.4 Analyze the Distribution of the Target Variable (Arrest Flag)

We will use crosstabs to summarize the relationship between each categorical feature and the “Arrest Flag” and bar plots to visualize the proportion of arrests within each category of the features.

```
[116]: # Create bar plots for each categorical feature against the "Arrest Flag"
plt.figure(figsize=(20, 20))
for i, col in enumerate(categorical_cols, 1):
    plt.subplot(3, 2, i)
    sns.countplot(x=col, hue='Arrest Flag', data=terry_df, order=terry_df[col].
    ↪value_counts().index)
    plt.title(f"Arrest Flag by {col}")
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[117]: # Print crosstab for each categorical feature with "Arrest Flag"
for col in categorical_cols:
    crosstab_result = pd.crosstab(terry_df[col], terry_df['Arrest Flag'])
    print(f"Crosstab for {col} with Arrest Flag:\n", crosstab_result, "\n")
```

Crosstab for Subject Age Group with Arrest Flag:

Subject Age Group	Arrest Flag	N	Y
-		2160	41
1 - 17		2155	131
18 - 25		10497	1085
26 - 35		18002	2387
36 - 45		11893	1735
46 - 55		6897	845
56 and Above		2856	366

Crosstab for Stop Resolution with Arrest Flag:

Stop Resolution	Arrest Flag	N	Y
Arrest		8210	6588

Citation / Infraction	219	0
Field Contact	29576	0
Offense Report	15729	0
Referred for Prosecution	726	2

Crosstab for Subject Perceived Race with Arrest Flag:

Arrest Flag	N	Y
Subject Perceived Race		
-	1718	98
American Indian or Alaska Native	1488	175
Asian	1815	273
Black or African American	16035	2286
Hispanic	1684	0
MULTIPLE SUBJECTS	2	0
Multi-Racial	809	0
Native Hawaiian or Other Pacific Islander	117	42
Other	152	0
Unknown	3851	569
White	26789	3147

Crosstab for Subject Perceived Gender with Arrest Flag:

Arrest Flag	N	Y
Subject Perceived Gender		
-	243	0
Female	11040	1100
Gender Diverse (gender non-conforming and/or tr...	35	10
MULTIPLE SUBJECTS	2	0
Male	42719	5475
Unable to Determine	326	0
Unknown	95	5

## 0.5 Statistical tests to validate the significance of observed patterns

We are using Chi-square Test for Independence

```
[118]: from scipy.stats import chi2_contingency

# Performing Chi-square test for 'Subject Perceived Race' and 'Arrest Flag'
contingency_table = pd.crosstab(terry_df['Subject Perceived Race'],
    ↪terry_df['Arrest Flag'])
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"Chi-square Test Results for 'Subject Perceived Race' and 'Arrest Flag':
    ↪\n Chi2: {chi2}, p-value: {p}")

# Interpretation of p-value
if p < 0.05:
```

```

    print("There is a significant association between 'Subject Perceived Race' and 'Arrest Flag'.")
else:
    print("No significant association between 'Subject Perceived Race' and 'Arrest Flag'.")

```

Chi-square Test Results for 'Subject Perceived Race' and 'Arrest Flag':  
Chi2: 503.15362641100324, p-value: 9.353668765101045e-102  
There is a significant association between 'Subject Perceived Race' and 'Arrest Flag'.

```

[119]: # Performing Chi-square test for 'Subject Perceived Gender' and 'Arrest Flag'
contingency_table = pd.crosstab(terry_df['Subject Perceived Gender'],
    ↪terry_df['Arrest Flag'])
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"Chi-square Test Results for 'Subject Perceived Gender' and 'Arrest Flag':\n Chi2: {chi2}, p-value: {p}")

# Interpretation of p-value
if p < 0.05:
    print("There is a significant association between 'Subject Perceived Gender' and 'Arrest Flag'.")
else:
    print("No significant association between 'Subject Perceived Gender' and 'Arrest Flag'.")

```

Chi-square Test Results for 'Subject Perceived Gender' and 'Arrest Flag':  
Chi2: 132.59728273046457, p-value: 3.647072311038753e-26  
There is a significant association between 'Subject Perceived Gender' and 'Arrest Flag'.

```

[120]: # Performing Chi-square test for 'Stop Resolution' and 'Arrest Flag'
contingency_table = pd.crosstab(terry_df['Stop Resolution'], terry_df['Arrest Flag'])
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"Chi-square Test Results for 'Stop Resolution' and 'Arrest Flag':\n Chi2: {chi2}, p-value: {p}")

# Interpretation of p-value
if p < 0.05:
    print("There is a significant association between 'Stop Resolution' and 'Arrest Flag'.")
else:
    print("No significant association between 'Stop Resolution' and 'Arrest Flag'.")

```



Chi-square Test Results for 'Stop Resolution' and 'Arrest Flag':

Chi2: 23071.40163433658, p-value: 0.0

There is a significant association between 'Stop Resolution' and 'Arrest Flag'.

```
[121]: # Performing Chi-square test for 'Subject Age Group' and 'Arrest Flag'
contingency_table = pd.crosstab(terry_df['Subject Age Group'], terry_df['Arrest_
    Flag'])
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"Chi-square Test Results for 'Subject Age Group' and 'Arrest Flag':\n
    Chi2: {chi2}, p-value: {p}")

# Interpretation of p-value
if p < 0.05:
    print("There is a significant association between 'Subject Age Group' and
    'Arrest Flag'.")
else:
    print("No significant association between 'Subject Age Group' and 'Arrest_
    Flag'.")
```

Chi-square Test Results for 'Subject Age Group' and 'Arrest Flag':

Chi2: 339.60818980749343, p-value: 2.6242186997924706e-70

There is a significant association between 'Subject Age Group' and 'Arrest Flag'.

### 0.5.1 Data Processing and modelling

We calculate the mode of the “Weapon Type” column and use it to fill any missing values in that column.

After filling the specified column, we drop rows with any remaining missing values in other columns.

```
[122]: # Drop other null columns
data_cleaned = terry_df.drop(columns=['Weapon Type', 'Officer Squad'])

data_cleaned.shape
```

[122]: (61050, 21)

```
[123]: # Check for missing values
data_cleaned.isnull().sum()
```

```
[123]: Subject Age Group      0
Subject ID                0
GO / SC Num              0
Terry Stop ID            0
Stop Resolution          0
Officer ID               0
```

```

Officer YOB          0
Officer Gender       0
Officer Race         0
Subject Perceived Race 0
Subject Perceived Gender 0
Reported Date        0
Reported Time        0
Initial Call Type    0
Final Call Type      0
Call Type            0
Arrest Flag          0
Frisk Flag           0
Precinct             0
Sector               0
Beat                0
dtype: int64

```

## 0.6 Encode Categorical Variables

I will use Label Encoding for simplicity, which converts each category to a unique integer.

```

[124]: from sklearn.preprocessing import LabelEncoder

# Categorical columns to encode
categorical_cols = ['Stop Resolution', 'Subject Perceived Race', 'Subject_
↳ Perceived Gender', 'Subject Age Group']

# Initialize label encoder
label_encoder = LabelEncoder()

# Encode categorical columns
for col in categorical_cols:
    data_cleaned[col] = label_encoder.fit_transform(data_cleaned[col]).
↳ astype(str))

data_cleaned.head()

```

```

[124]:
  Subject Age Group  Subject ID  GO / SC Num  Terry Stop ID  \
0                1         -1  20170000315550        301338
1                3         -1  20170000115391        253596
2                2         -1  20160000036392        123863
3                4  33919122751  20210000055627  21770133324
4                2         -1  20160000003522        183376

  Stop Resolution  Officer ID  Officer YOB  Officer Gender  Officer Race  \
0                3        5489        1964              M        White
1                3        6403        1969              M        White

```

2	3	7473	1981	M	White
3	2	6711	1977	M	White
4	2	7090	1981	F	White

	Subject Perceived Race	...	Reported Date	Reported Time	\
0	10	...	2017-08-26T00:00:00Z	00:10:00.0000000	
1	10	...	2017-04-05T00:00:00Z	16:04:00.0000000	
2	10	...	2016-01-31T00:00:00Z	02:26:00.0000000	
3	10	...	2021-03-06T00:00:00Z	14:42:43.0000000	
4	10	...	2016-08-14T00:00:00Z	15:49:00.0000000	

	Initial Call Type	\
0	NUISANCE	
1	THEFT (DOES NOT INCLUDE SHOPLIFT OR SVCS)	
2	SUSPICIOUS PERSON, VEHICLE, OR INCIDENT	
3	PROPERTY - DAMAGE	
4	-	

	Final Call Type	Call Type	\
0	--MISCHIEF OR NUISANCE - GENERAL	911	
1	--THEFT - ALL OTHER TELEPHONE OTHER, NOT	911	
2	--SUSPICIOUS CIRCUM. - SUSPICIOUS PERSON	911	
3	--DISTURBANCE - OTHER	911	
4	-	-	

	Arrest Flag	Frisk Flag	Precinct	Sector	Beat
0	N	N	North	N	N2
1	N	N	North	L	L3
2	N	N	North	B	B3
3	N	N	-	-	-
4	N	Y	-	-	-

[5 rows x 21 columns]

## 0.7 Feature Scaling and Splitting the data

```
[125]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

# Identify numerical columns
numerical_cols = data_cleaned.select_dtypes(include=[np.number]).columns.
               ↪ tolist()
numerical_cols
```

```
[125]: ['Subject Age Group',
        'Subject ID',
        'GO / SC Num',
```

```
'Terry Stop ID',
'Stop Resolution',
'Officer YOB',
'Subject Perceived Race',
'Subject Perceived Gender']
```

```
[126]: scaler = StandardScaler()

# Scale numerical columns
data_cleaned[numerical_cols] = scaler.
    ↪ fit_transform(data_cleaned[numerical_cols])

data_cleaned.head()
```

```
[126]: Subject Age Group Subject ID GO / SC Num Terry Stop ID Stop Resolution \
0          -1.686045   -0.572752   -0.194557   -0.694398         1.085697
1          -0.195569   -0.572752   -0.194560   -0.694400         1.085697
2          -0.940807   -0.572752   -0.311228   -0.694408         1.085697
3           0.549670    2.090833    0.272110    0.544906         0.186107
4          -0.940807   -0.572752   -0.311229   -0.694404         0.186107
```

```
Officer ID Officer YOB Officer Gender Officer Race Subject Perceived Race \
0         5489   -2.120560                M      White         0.892841
1         6403   -1.592669                M      White         0.892841
2         7473   -0.325732                M      White         0.892841
3         6711   -0.748045                M      White         0.892841
4         7090   -0.325732                F      White         0.892841
```

```
... Reported Date Reported Time \
0 ... 2017-08-26T00:00:00Z 00:10:00.0000000
1 ... 2017-04-05T00:00:00Z 16:04:00.0000000
2 ... 2016-01-31T00:00:00Z 02:26:00.0000000
3 ... 2021-03-06T00:00:00Z 14:42:43.0000000
4 ... 2016-08-14T00:00:00Z 15:49:00.0000000
```

```
Initial Call Type \
0 NUISANCE
1 THEFT (DOES NOT INCLUDE SHOPLIFT OR SVCS)
2 SUSPICIOUS PERSON, VEHICLE, OR INCIDENT
3 PROPERTY - DAMAGE
4 -
```

```
Final Call Type Call Type \
0 --MISCHIEF OR NUISANCE - GENERAL 911
1 --THEFT - ALL OTHER TELEPHONE OTHER, NOT 911
2 --SUSPICIOUS CIRCUM. - SUSPICIOUS PERSON 911
3 --DISTURBANCE - OTHER 911
```

4

-

-

	Arrest	Flag	Frisk	Flag	Precinct	Sector	Beat
0		N		N	North	N	N2
1		N		N	North	L	L3
2		N		N	North	B	B3
3		N		N	-	-	-
4		N		Y	-	-	-

[5 rows x 21 columns]

```
[127]: # Split the data into features and target variable
X = data_cleaned[categorical_cols]
y = data_cleaned['Arrest Flag']
```

```
[128]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

print("Training set size:", X_train.shape)
print("Testing set size:", X_test.shape)
```

Training set size: (48840, 4)

Testing set size: (12210, 4)

## 0.8 Model Building

```
[129]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix,
    accuracy_score

# Initialize models
logistic_regression = LogisticRegression()
decision_tree = DecisionTreeClassifier()

# Train Logistic Regression model
logistic_regression.fit(X_train, y_train)

# Train Decision Tree model
decision_tree.fit(X_train, y_train)
```

[129]: DecisionTreeClassifier()

```
[130]: # Predict on the test set with Logistic Regression
y_pred_lr = logistic_regression.predict(X_test)
```

```
[131]: # Predict on the test set with Decision Tree
y_pred_dt = decision_tree.predict(X_test)
```

```
[132]: # Evaluate Logistic Regression model
print("Logistic Regression Model Evaluation:")
print("Accuracy:", accuracy_score(y_test, y_pred_lr))
print("Classification Report:\n", classification_report(y_test, y_pred_lr))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_lr))
```

Logistic Regression Model Evaluation:

Accuracy: 0.8925470925470925

Classification Report:

	precision	recall	f1-score	support
N	0.91	0.98	0.94	10921
Y	0.47	0.16	0.24	1289
accuracy			0.89	12210
macro avg	0.69	0.57	0.59	12210
weighted avg	0.86	0.89	0.87	12210

Confusion Matrix:

```
[[10696  225]
 [ 1087  202]]
```

```
[133]: # Evaluate Decision Tree model
print("\nDecision Tree Model Evaluation:")
print("Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Classification Report:\n", classification_report(y_test, y_pred_dt))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
```

Decision Tree Model Evaluation:

Accuracy: 0.9013923013923014

Classification Report:

	precision	recall	f1-score	support
N	0.93	0.96	0.95	10921
Y	0.55	0.40	0.46	1289
accuracy			0.90	12210
macro avg	0.74	0.68	0.70	12210
weighted avg	0.89	0.90	0.89	12210

Confusion Matrix:

```
[[10492  429]
 [  775  514]]
```

## 0.9 Hyperparameter Tuning

```
[134]: # Import GridSearch library
from sklearn.model_selection import GridSearchCV

[135]: # Iterative Development: Grid Search for Hyperparameter Tuning for Decision Tree
param_grid = {'max_depth': [None, 10, 20, 30], 'min_samples_split': [2, 5, 10]}
grid_search_dt = GridSearchCV(decision_tree, param_grid, cv=5)
grid_search_dt.fit(X_train, y_train)

[135]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                  param_grid={'max_depth': [None, 10, 20, 30],
                              'min_samples_split': [2, 5, 10]})

[136]: # Best model from Grid Search
best_dt = grid_search_dt.best_estimator_

[137]: # Predict using the best Decision Tree model
y_pred_best_dt = best_dt.predict(X_test)

[138]: # Evaluate the best Decision Tree model
print("\nBest Decision Tree Model Evaluation after Grid Search:")
print("Accuracy:", accuracy_score(y_test, y_pred_best_dt))
print("Classification Report:\n", classification_report(y_test, y_pred_best_dt))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_best_dt))
```

Best Decision Tree Model Evaluation after Grid Search:

Accuracy: 0.9013104013104013

Classification Report:

	precision	recall	f1-score	support
N	0.93	0.96	0.95	10921
Y	0.54	0.40	0.46	1289
accuracy			0.90	12210
macro avg	0.74	0.68	0.70	12210
weighted avg	0.89	0.90	0.89	12210

Confusion Matrix:


```
[[10491  430]
 [ 775  514]]
```

## 0.10 Visualizations

```
[139]: # Import libraries
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, roc_auc_score
```

### 0.10.1 Confusion Matrix for both models


```
[140]: # Confusion Matrix for Logistic Regression
cm_lr = confusion_matrix(y_test, y_pred_lr)
disp_lr = ConfusionMatrixDisplay(confusion_matrix=cm_lr,
    ↳display_labels=logistic_regression.classes_)
disp_lr.plot(cmap='Blues')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()
```



```
[141]: # Confusion Matrix for Decision Tree
cm_dt = confusion_matrix(y_test, y_pred_dt)
disp_dt = ConfusionMatrixDisplay(confusion_matrix=cm_dt,
    ↳display_labels=decision_tree.classes_)
```

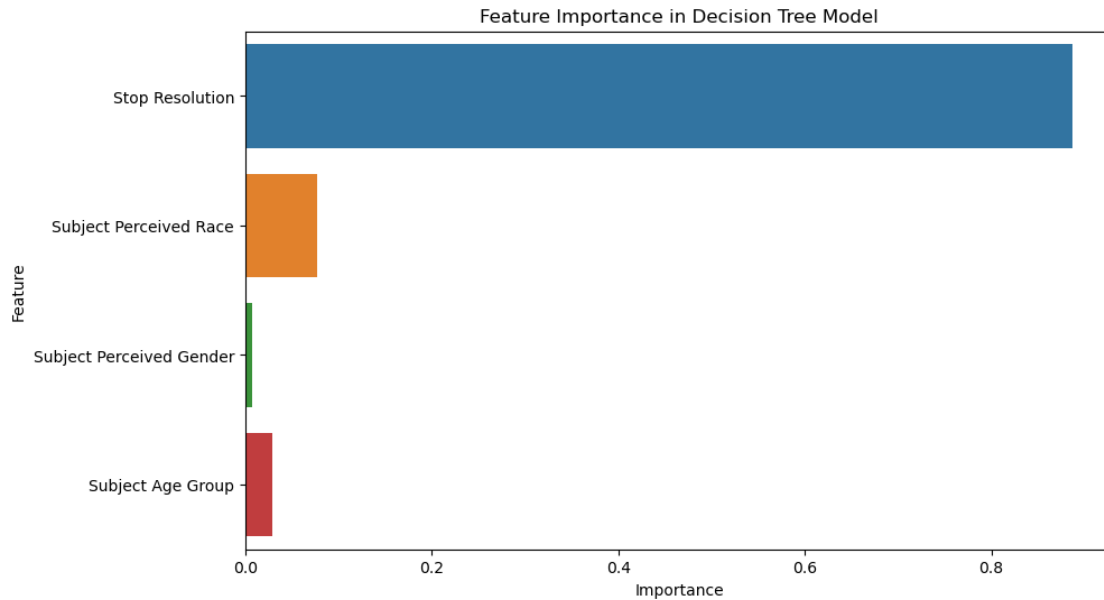


```
disp_dt.plot(cmap='Greens')
plt.title('Confusion Matrix for Decision Tree')
plt.show()
```



### 0.10.2 Feature Importance for Decision Tree


```
[142]: # Feature Importance for Decision Tree
plt.figure(figsize=(10, 6))
sns.barplot(x=decision_tree.feature_importances_, y=categorical_cols)
plt.title('Feature Importance in Decision Tree Model')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```



### 0.10.3 Coefficients for Logistic Regression

```
[143]: # Coefficients for Logistic Regression
coefficients = pd.DataFrame({'Feature': X_train.columns, 'Coefficient': logistic_regression.coef_[0]})
coefficients = coefficients.sort_values(by='Coefficient', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='Coefficient', y='Feature', data=coefficients)
plt.title('Feature Coefficients in Logistic Regression Model')
plt.xlabel('Coefficient')
plt.ylabel('Feature')
plt.show()
```



#### 0.10.4 ROC Curve

```
[144]: # Convert 'No' to 0 and 'Yes' to 1 in y_test and predictions
y_test_binary = y_test.map({'N': 0, 'Y': 1})
y_pred_prob_lr = logistic_regression.predict_proba(X_test)[: , 1]
y_pred_prob_dt = decision_tree.predict_proba(X_test)[: , 1]
```

```
[145]: # Prediction probabilities for Logistic Regression
y_pred_prob_lr = logistic_regression.predict_proba(X_test)[: , 1]

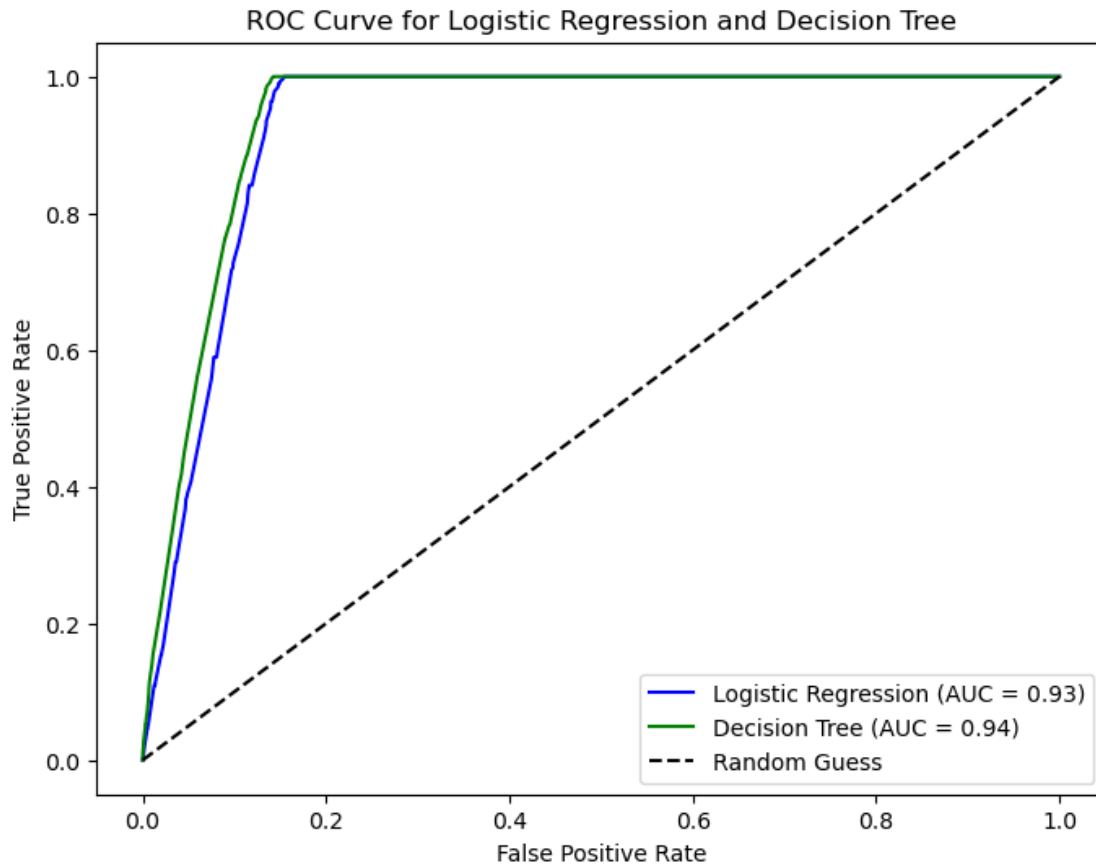
# Prediction probabilities for Decision Tree
y_pred_prob_dt = decision_tree.predict_proba(X_test)[: , 1]

# ROC Curve for Logistic Regression
fpr_lr, tpr_lr, _ = roc_curve(y_test_binary, y_pred_prob_lr)
roc_auc_lr = roc_auc_score(y_test_binary, y_pred_prob_lr)

# ROC Curve for Decision Tree
fpr_dt, tpr_dt, _ = roc_curve(y_test_binary, y_pred_prob_dt)
roc_auc_dt = roc_auc_score(y_test_binary, y_pred_prob_dt)

# Plotting the ROC curves on a single graph
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label=f'Logistic Regression (AUC = {roc_auc_lr:.2f})',
        color='blue')
plt.plot(fpr_dt, tpr_dt, label=f'Decision Tree (AUC = {roc_auc_dt:.2f})',
        color='green')
```

```
plt.plot([0, 1], [0, 1], 'k--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Logistic Regression and Decision Tree')
plt.legend(loc='lower right')
plt.show()
```



## 0.11 Cross-Validation

```
[146]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, accuracy_score, precision_score,
      ↪ recall_score, f1_score
```

```
[147]: # Define scoring metrics
scoring = {'accuracy': make_scorer(accuracy_score),
           'precision': make_scorer(precision_score, pos_label='Y'),
           'recall': make_scorer(recall_score, pos_label='Y'),
           'f1': make_scorer(f1_score, pos_label='Y')}
```

```
[148]: # Cross-validation for Logistic Regression
cv_scores_lr = cross_val_score(logistic_regression, X_train, y_train, cv=5,
    ↳scoring='accuracy')
print(f"Cross-Validation Scores for Logistic Regression (Accuracy):\n
    ↳{cv_scores_lr}")
print(f"Average Cross-Validation Score for Logistic Regression (Accuracy):\n
    ↳{cv_scores_lr.mean()}")

# If you want to calculate multiple metrics at once:
from sklearn.model_selection import cross_validate
cv_results_lr = cross_validate(logistic_regression, X_train, y_train, cv=5,
    ↳scoring=scoring)

# Print detailed cross-validation results
print("\nDetailed Cross-Validation Results for Logistic Regression:")
for key in scoring.keys():
    print(f"{key.capitalize()} Scores: {cv_results_lr['test_' + key]}")
    print(f"Average {key.capitalize()} Score: {cv_results_lr['test_' + key].
    ↳mean()}")
```

Cross-Validation Scores for Logistic Regression (Accuracy): [0.89209664  
0.89045864 0.89250614 0.89322277 0.88841114]  
Average Cross-Validation Score for Logistic Regression (Accuracy):  
0.8913390663390663

Detailed Cross-Validation Results for Logistic Regression:  
Accuracy Scores: [0.89209664 0.89045864 0.89250614 0.89322277 0.88841114]  
Average Accuracy Score: 0.8913390663390663  
Precision Scores: [0.505 0.4872449 0.50853242 0.52449568 0.45797101]  
Average Precision Score: 0.4966488025787116  
Recall Scores: [0.28584906 0.18018868 0.28113208 0.17169811 0.14891612]  
Average Recall Score: 0.21355680827983567  
F1 Scores: [0.36506024 0.2630854 0.36208991 0.25870647 0.22475107]  
Average F1 Score: 0.2947386179752469

```
[149]: # Cross-validation for Decision Tree
cv_scores_dt = cross_val_score(decision_tree, X_train, y_train, cv=5,
    ↳scoring='accuracy')
print(f"\nCross-Validation Scores for Decision Tree (Accuracy): {cv_scores_dt}")
print(f"Average Cross-Validation Score for Decision Tree (Accuracy):\n
    ↳{cv_scores_dt.mean()}")

# Calculate multiple metrics at once:
cv_results_dt = cross_validate(decision_tree, X_train, y_train, cv=5,
    ↳scoring=scoring)

# Print detailed cross-validation results
```

```

print("\nDetailed Cross-Validation Results for Decision Tree:")
for key in scoring.keys():
    print(f"{key.capitalize()} Scores: {cv_results_dt['test_' + key]}")
    print(f"Average {key.capitalize()} Score: {cv_results_dt['test_' + key].
↪mean()}")

```

Cross-Validation Scores for Decision Tree (Accuracy): [0.90120803 0.89946765  
0.8996724 0.90284603 0.90120803]  
Average Cross-Validation Score for Decision Tree (Accuracy): 0.9008804258804259

Detailed Cross-Validation Results for Decision Tree:  
Accuracy Scores: [0.90120803 0.89946765 0.8996724 0.90284603 0.90120803]  
Average Accuracy Score: 0.9008804258804259  
Precision Scores: [0.56129032 0.54779412 0.55899705 0.58295964 0.56282723]  
Average Precision Score: 0.5627736713523384  
Recall Scores: [0.41037736 0.42169811 0.35754717 0.36792453 0.40527804]  
Average Recall Score: 0.3925650418793235  
F1 Scores: [0.47411444 0.47654584 0.43613349 0.45112782 0.47123288]  
Average F1 Score: 0.4618308933323954